

Chanhee Park

📍 Seoul, Republic of Korea ✉ pch7678@korea.ac.kr

🌐 johnnynlp.github.io 🏠 github.com/JohnnyNLP 🎓 Google Scholar 🔗 linkedin.com/in/chanhee-park-nlp

EXPERIENCE

Tech Lead, AI Engineer (freelance)

Oct. 2025 – Present

LOUM Partner: LLM platform for HR and labor-law workflows

Seoul, Korea

- Developed a labor-law expert agent that answers legal questions, drafts legal documents, and reviews compliance through a single entry point; optimized its **Planner / Executor / Composer** architecture down to **20 s** response time on document-heavy requests
- Built retrieval over 10 public Korean legal databases with 127K document chunks, web search, and attorney-curated documents on legal-fine-tuned KURE embeddings; shipped document generation, review, and consulting on AWS. **1.1K users** through closed beta
- Built the internal evaluation harness for the product's 6 task types with over 1,000 practitioner-sourced queries, A/B ablations, and model-migration gates, all monitored through a self-maintained debug portal with per-request cost and token tracing

Interpreting Officer (First Lieutenant)

Mar. 2020 – Jun. 2023

Republic of Korea Air Force, Logistics Command

Daegu, Korea

- Lead interpreter for defense procurement with Boeing, BAE Systems, and Lockheed Martin, including quadrilateral ROKAF-USAF-DAPA talks at Boeing St. Louis; ROK-US Combined Exercises; Black Eagles airshow tour across 4 countries

FUNDED PROJECTS

Sovereign AI Foundation Model | NC AI

Sep. – Dec. 2025

- SFT cold-start data for **VAETKI-100B-A10B**: 60+ open datasets plus AIHub and in-house sources across 4 languages and 9 domains, curated to **6.5M Q/A samples** via 6-dimension clarity scoring and consistency-triggered human review

LegalMi:dm | KT

Jul. 2024 – Jul. 2025

- Legal-domain continual pretraining on Mi:dm-2.0-Base (11.5B) with DeepSpeed multi-node training; the model was central to KT's successful bid in a **Supreme Court of Korea** consortium

Clinical RAG for Breast Cancer Decision Support | Korea University

May – Dec. 2024

- Led end-to-end as principal researcher: 1,356 ASCO guideline documents, 98 clinician queries, FAISS + SentenceTransformers retrieval; showed small open-source models with RAG can approach GPT-4 on clinical decision support

Korean Writing & Grammar Assistant | KOCCA & NIKL

Aug. 2026 – Present

- Designed the core evaluation metrics used to officially certify the agent across five tasks

FIRST-AUTHORED PUBLICATIONS

Can Small Open-Source Language Models With RAG Match GPT-4 Performance in Breast Cancer Clinical Decision Support? C. Park, I. H. Park, M. Kim, Y. Jang, H. Lim, Y. Hur. *JCO Clinical Cancer Informatics*, 2026

MIRAGE: A Metric-Intensive Benchmark for Retrieval-Augmented Generation Evaluation. C. Park, H. Moon, H. Lim. *Findings of NAACL 2025*

AgentHop: A Diagnostic Benchmark for Agentic Multi-Hop Scientific Question Answering. C. Park, J. Yoon, S. Han, H. Moon, H. Lim. *NeurIPS Evaluations & Datasets*, 2026 (under review)

Funnel of Thoughts: Efficient Test-Time Scaling via Early Voting and Rollout Pruning. C. Park, S. Han, J. Yoon, S. Hong, H. Lim. *AACL-IJCNLP 2026* (under review)

Ko-MIRAGE: A Diagnostic Benchmark for Korean Retrieval-Augmented Generation. C. Park, Y. Hur, H. Lim. *HCLT 2025*

DOWA: Document-Augmented Word Assessment through Dual-Labeling Approach. C. Park, C. Park, A. I. Dainelles, H. Lim. *HCLT 2024*

CO-AUTHORED PUBLICATIONS

Towards Privacy-Preserving Large Language Model: Text-free Inference Through Alignment and Adaptation. J. Yoon, **C. Park**, Y. Chun, H. Moon, H. Lim. *ACL 2026 Main*

Evidential Transformation Network: Turning Pretrained Models into Evidential Models for Post-hoc Uncertainty Estimation. Y. Chun, **C. Park**, J. Yoon, J. Seo, H. Lim. *CVPR 2026*

ZENA: Enhanced Anomaly-Based IDS for Smart Homes Automation through Multilayer Artificial Neural Network. N. K. Carlos, **C. Park**, I. Aiyanyo, H. Lee, H. Lim. *Journal of Communications and Networks, 2026*

LegalMi:dm: Use-Case-Driven Legal Domain Specialization for Korean Large Language Model. Y. Jang, **C. Park**, H. Moon, Y. Ham, J. Moon, J. Kim, J. Jung, H. Lim. *DATA-FM Workshop @ ICLR 2026*

SchemaPrism: A Diagnostic Benchmark for Profiling LLMs' Schema Linking Capability in Text-to-SQL. M. Kang, J. Son, **C. Park**, E. H. Prince, H. Lim. *AAAI 2026* (under review)

KITE: A Benchmark for Evaluating Korean Instruction-Following Abilities in Large Language Models. D. Kim, **C. Park**, H. Lim. *arXiv preprint, 2025*

Database Retrieval-Augmented Generation Framework Leveraging Large Reasoning Language Models and Schema Linking Methods. M. Kang, J. Son, **C. Park**, A. So, H. Lim. *HCLT 2025*

VAETKI Technical Report. NC-AI Consortium, data team contributor. *Technical report, 2026*

PATENTS

Device and Method for Evaluating Retrieval-Augmented Generation Systems (MIRAGE). Filed KR, Dec. 2025

Method and Apparatus for Accelerating Multi-Sample Inference of Large Reasoning Model (FoT). Filed KR, Aug. 2026

EDUCATION

Korea University <i>M.S., Computer Science and Engineering</i>	Seoul, Korea <i>Feb. 2025 – Feb. 2027 (expected)</i>
Goorm AI Natural Language Processing Course (12th) <i>16-week intensive course work on linear algebra, statistics, machine learning, and Python</i>	Seoul, Korea <i>Jun. – Oct. 2023</i>
Korea University <i>B.A., Korean Language and Literature; second major, German Literature</i>	Seoul, Korea <i>Mar. 2014 – Feb. 2020</i>

TECHNICAL SKILLS

Serving: vLLM, SGLang, Amazon SageMaker, sentence-transformers

Training: PyTorch, Transformers, DeepSpeed, TRL, Ray

Agent & RAG: FastAPI, SSE streaming, LangChain/LangGraph/LangSmith, multi-provider APIs (Gemini, Claude, OpenAI, Together, DeepSeek), FAISS, MongoDB

Data & Evaluation: datatrove, lm-evaluation-harness, evalchemy, FactScore, MTEB

Infrastructure: CUDA, Python, AWS (EC2, S3, Lambda, CloudWatch), Docker, Tailscale

Languages: Korean (native), English (professional, TOEIC 980), German (professional, TestDaF 4/5/4/5)